

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
%matplotlib inline
```

```
from __future__ import print_function
%matplotlib inline
import argparse
import os
import random
import torch
import torch.nn as nn
import torch.nn.parallel
import torch.backends.cudnn as cudnn
import torch.optim as optim
import torch.utils.data
import torchvision.datasets as dset
import torchvision.transforms as transforms
import torchvision.utils as vutils
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.animation as animation
from IPython.display import HTML
import zipfile
import cv2

# Set random seem for reproducibility
manualSeed = 999
#manualSeed = random.randint(1, 10000) # use if you want new results
print("Random Seed: ", manualSeed)
random.seed(manualSeed)
torch.manual_seed(manualSeed)
```

Random Seed: 999
<torch._C.Generator at 0x7f403a2e52d0>

```
drive = '/content/sample_data/square_data/squares/'
for i in range(5000):
    x = random.randrange(0,10)
    img = np.zeros((64,64,3), np.uint8)
    color = (random.randrange(1,255), random.randrange(1,255), random.randrange(1,255))

    img[32-x:32+x, 32-x:32+x, :] = tuple(reversed(color))
    rotationMat = cv2.getRotationMatrix2D((32,32), random.randrange(0,90), 1)
    img = cv2.warpAffine(img, rotationMat,(img.shape[0],img.shape[1]))
    trans_mat = np.float32([[1,0,random.randrange(-20,20)], [0,1,random.randrange(-20,20)]])
    img = cv2.warpAffine(img, trans_mat, (img.shape[0],img.shape[1]))
    # cv2.war
    file = drive + 'square'+ str(i) + '.png'
    plt.imsave(file, img)
```

▼ Inputs

Let's define some inputs for the run:

- **dataroot** - the path to the root of the dataset folder. We will talk more about the dataset in the next section
- **workers** - the number of worker threads for loading the data with the DataLoader
- **batch_size** - the batch size used in training. The DCGAN paper uses a batch size of 128
- **image_size** - the spatial size of the images used for training. This implementation defaults to 64. If a larger size is desired, the structures of D and G must be changed. See [here](https://github.com/pytorch/examples/issues/70) for more details
- **nc** - number of color channels in the input images. For color images this is 3
- **nz** - length of latent vector
- **ngf** - relates to the depth of feature maps carried through the generator
- **ndf** - sets the depth of feature maps propagated through the discriminator
- **num_epochs** - number of training epochs to run. Training for longer will probably lead to better results but also take much longer
- **lr** - learning rate for training. As described in the DCGAN paper, this number should be 0.0002
- **beta1** - beta1 hyperparameter for Adam optimizers. As described in paper, this number should be 0.5
- **ngpu** - number of GPUs available. If this is 0, code will run in CPU mode. If this number is greater than 0, code will run on that number of GPUs

```
# Root directory for dataset
dataroot = "/content/sample_data/square_data"

# Number of workers for dataloader
workers = 4

# Batch size during training
batch_size = 64

# Spatial size of training images. All images will be resized to this
# size using a transformer.
image_size = 64

# Number of channels in the training images. For color images this is 3
nc = 3

# Size of z latent vector (i.e. size of generator input)
nz = 100

# Size of feature maps in generator
ngf = 64

# Size of feature maps in discriminator
ndf = 64

# Number of training epochs
num_epochs = 60

# Learning rate for optimizers
lr = 0.001

# Beta1 hyperparam for Adam optimizers
```

```
beta1 = 0.5
```

```
# Number of GPUs available. Use 0 for CPU mode.
```

```
ngpu = 1
```

```
# We can use an image folder dataset the way we have it setup.
```

```
# Create the dataset
```

```
dataset = dset.ImageFolder(root=dataroot,
                           transform=transforms.Compose([
                               transforms.Resize(image_size),
                               transforms.CenterCrop(image_size),
                               transforms.ToTensor(),
                               transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
                           ]))
```

```
# Create the dataloader
```

```
dataloader = torch.utils.data.DataLoader(dataset, batch_size=batch_size,
                                          shuffle=True, num_workers=workers)
```

```
# Decide which device we want to run on
```

```
device = torch.device("cuda:0" if (torch.cuda.is_available() and ngpu > 0) else "cpu")
```

```
# Plot some training images
```

```
real_batch = next(iter(dataloader))
```

```
plt.figure(figsize=(8,8))
```

```
plt.axis("off")
```

```
plt.title("Training Images")
```

```
plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:64], padding=2, normalize=True)).cpu(),
```

☞ <matplotlib.image.AxesImage at 0x7f4084e20898>

Training Images



```

# custom weights initialization called on netG and netD
def weights_init(m):
    classname = m.__class__.__name__
    if classname.find('Conv') != -1:
        nn.init.normal_(m.weight.data, 0.0, 0.02)
    elif classname.find('BatchNorm') != -1:
        nn.init.normal_(m.weight.data, 1.0, 0.02)
        nn.init.constant_(m.bias.data, 0)

# Generator Code

class Generator(nn.Module):
    def __init__(self, ngpu):
        super(Generator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is Z, going into a convolution
            nn.ConvTranspose2d( nz, ngf * 8, 4, 1, 0, bias=False),
            nn.BatchNorm2d(ngf * 8),
            nn.ReLU(True),
            # state size. (ngf*8) x 4 x 4
            nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True),
            # state size. (ngf*4) x 8 x 8
            nn.ConvTranspose2d( ngf * 4, ngf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 2),
            nn.ReLU(True),
            # state size. (ngf*2) x 16 x 16
            nn.ConvTranspose2d( ngf * 2, ngf, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf),
            nn.ReLU(True),
            # state size. (ngf) x 32 x 32
            nn.ConvTranspose2d( ngf, nc, 4, 2, 1, bias=False),
            nn.Tanh()
            # state size. (nc) x 64 x 64
        )

    def forward(self, input):
        return self.main(input)

```

Now, we can instantiate the generator and apply the `weights_init` function. Check out the printed `m` the generator object is structured.

```

# Create the generator
netG = Generator(ngpu).to(device)

# Handle multi-gpu if desired
if (device.type == 'cuda') and (ngpu > 1):
    netG = nn.DataParallel(netG, list(range(ngpu)))

# Apply the weights_init function to randomly initialize all weights
# to mean=0, stdev=0.2.
netG.apply(weights_init)

# Print the model
print(netG)

```



```

Generator(
  (main): Sequential(
    (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU(inplace=True)
    (9): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (11): ReLU(inplace=True)
    (12): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (13): Tanh()
  )
)

```

Discriminator Code

```

class Discriminator(nn.Module):
    def __init__(self, ngpu):
        super(Discriminator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is (nc) x 64 x 64
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf) x 32 x 32
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 2),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*2) x 16 x 16
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 4),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*4) x 8 x 8
            nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 8),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*8) x 4 x 4
            nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),
            nn.Sigmoid()
        )

    def forward(self, input):
        return self.main(input)

```

Now, as with the generator, we can create the discriminator, apply the `weights_init` function, and print its structure.

```

# Create the Discriminator
netD = Discriminator(ngpu).to(device)

# Handle multi-gpu if desired

```

```

if (device.type == 'cuda') and (ngpu > 1):
    netD = nn.DataParallel(netD, list(range(ngpu)))

# Apply the weights_init function to randomly initialize all weights
# to mean=0, stdev=0.2.
netD.apply(weights_init)

# Print the model
print(netD)

[ ]> Discriminator(
  (main): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (10): LeakyReLU(negative_slope=0.2, inplace=True)
    (11): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (12): Sigmoid()
  )
)

# Initialize BCELoss function
criterion = nn.BCELoss()

# Create batch of latent vectors that we will use to visualize
# the progression of the generator
fixed_noise = torch.randn(64, nz, 1, 1, device=device)

# Establish convention for real and fake labels during training
real_label = 1
fake_label = 0

# Setup Adam optimizers for both G and D
optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))

# Training Loop

# Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
iters = 0

print("Starting Training Loop...")
# For each epoch
for epoch in range(num_epochs):
    # For each batch in the dataloader
    for i, data in enumerate(dataloader, 0):

        #####
        # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
        #####
        ## Train with all-real batch

```

```

netD.zero_grad()
# Format batch
real_cpu = data[0].to(device)
b_size = real_cpu.size(0)
label = torch.full((b_size,), real_label, device=device)
# Forward pass real batch through D
output = netD(real_cpu).view(-1)
# Calculate loss on all-real batch
errD_real = criterion(output, label)
# Calculate gradients for D in backward pass
errD_real.backward()
D_x = output.mean().item()

## Train with all-fake batch
# Generate batch of latent vectors
noise = torch.randn(b_size, nz, 1, 1, device=device)
# Generate fake image batch with G
fake = netG(noise)
label.fill_(fake_label)
# Classify all fake batch with D
output = netD(fake.detach()).view(-1)
# Calculate D's loss on the all-fake batch
errD_fake = criterion(output, label)
# Calculate the gradients for this batch
errD_fake.backward()
D_G_z1 = output.mean().item()
# Add the gradients from the all-real and all-fake batches
errD = errD_real + errD_fake
# Update D
optimizerD.step()

#####
# (2) Update G network: maximize log(D(G(z)))
#####
netG.zero_grad()
label.fill_(real_label) # fake labels are real for generator cost
# Since we just updated D, perform another forward pass of all-fake batch through D
output = netD(fake).view(-1)
# Calculate G's loss based on this output
errG = criterion(output, label)
# Calculate gradients for G
errG.backward()
D_G_z2 = output.mean().item()
# Update G
optimizerG.step()

# Output training stats
if i % 50 == 0:
    print('[%d/%d][%d/%d]\tLoss_D: %.4f\tLoss_G: %.4f\tD(x): %.4f\tD(G(z)): %.4f / %.4f'
          % (epoch, num_epochs, i, len(dataloader),
             errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))

# Save Losses for plotting later
G_losses.append(errG.item())
D_losses.append(errD.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i == len(dataloader)-1)):
    with torch.no_grad():
        fake = netG(fixed_noise).detach().cpu()
        img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

iters += 1

```



[20/60][50/79]	Loss_D: 1.4951	Loss_G: 0.7221	D(x): 0.3653	D(G(z)): 0.2892 / 0.5438
[21/60][0/79]	Loss_D: 2.0199	Loss_G: 1.5116	D(x): 0.9083	D(G(z)): 0.7975 / 0.2521
[21/60][50/79]	Loss_D: 1.2511	Loss_G: 1.4097	D(x): 0.8450	D(G(z)): 0.6402 / 0.2637
[22/60][0/79]	Loss_D: 2.0815	Loss_G: 2.3464	D(x): 0.8709	D(G(z)): 0.8087 / 0.1175
[22/60][50/79]	Loss_D: 1.2457	Loss_G: 1.5257	D(x): 0.6614	D(G(z)): 0.5218 / 0.2515
[23/60][0/79]	Loss_D: 1.4509	Loss_G: 1.4848	D(x): 0.8800	D(G(z)): 0.6792 / 0.2593
[23/60][50/79]	Loss_D: 0.9505	Loss_G: 1.1233	D(x): 0.5841	D(G(z)): 0.3002 / 0.3478
[24/60][0/79]	Loss_D: 1.5079	Loss_G: 1.7092	D(x): 0.9072	D(G(z)): 0.7195 / 0.2453
[24/60][50/79]	Loss_D: 1.1446	Loss_G: 0.6574	D(x): 0.4374	D(G(z)): 0.2374 / 0.5424
[25/60][0/79]	Loss_D: 2.7688	Loss_G: 0.9737	D(x): 0.9293	D(G(z)): 0.8896 / 0.4313
[25/60][50/79]	Loss_D: 1.2217	Loss_G: 1.3049	D(x): 0.4208	D(G(z)): 0.2370 / 0.2869
[26/60][0/79]	Loss_D: 3.3645	Loss_G: 0.6042	D(x): 0.9731	D(G(z)): 0.9400 / 0.6066
[26/60][50/79]	Loss_D: 1.2905	Loss_G: 1.1599	D(x): 0.5157	D(G(z)): 0.4129 / 0.3615
[27/60][0/79]	Loss_D: 1.5698	Loss_G: 1.4661	D(x): 0.8260	D(G(z)): 0.6743 / 0.2817
[27/60][50/79]	Loss_D: 1.3573	Loss_G: 2.9691	D(x): 0.8115	D(G(z)): 0.6475 / 0.0616
[28/60][0/79]	Loss_D: 2.6460	Loss_G: 1.8834	D(x): 0.9850	D(G(z)): 0.8910 / 0.1909
[28/60][50/79]	Loss_D: 1.6447	Loss_G: 3.1307	D(x): 0.9376	D(G(z)): 0.7692 / 0.0492
[29/60][0/79]	Loss_D: 1.6191	Loss_G: 2.7252	D(x): 0.7974	D(G(z)): 0.6507 / 0.1024
[29/60][50/79]	Loss_D: 1.3014	Loss_G: 0.9880	D(x): 0.3335	D(G(z)): 0.0545 / 0.4108
[30/60][0/79]	Loss_D: 0.8916	Loss_G: 4.9544	D(x): 0.9233	D(G(z)): 0.5092 / 0.0087
[30/60][50/79]	Loss_D: 0.4139	Loss_G: 2.6281	D(x): 0.8747	D(G(z)): 0.2275 / 0.0873
[31/60][0/79]	Loss_D: 1.7637	Loss_G: 2.2514	D(x): 0.8812	D(G(z)): 0.7436 / 0.1495
[31/60][50/79]	Loss_D: 0.4079	Loss_G: 2.4651	D(x): 0.7282	D(G(z)): 0.0678 / 0.1001
[32/60][0/79]	Loss_D: 0.6146	Loss_G: 0.7553	D(x): 0.6417	D(G(z)): 0.1217 / 0.5311
[32/60][50/79]	Loss_D: 0.1896	Loss_G: 3.1153	D(x): 0.9523	D(G(z)): 0.1264 / 0.0528
[33/60][0/79]	Loss_D: 1.6854	Loss_G: 3.0538	D(x): 0.8678	D(G(z)): 0.7134 / 0.0717
[33/60][50/79]	Loss_D: 0.5339	Loss_G: 5.1051	D(x): 0.9929	D(G(z)): 0.3859 / 0.0072
[34/60][0/79]	Loss_D: 0.4652	Loss_G: 3.1988	D(x): 0.8242	D(G(z)): 0.1858 / 0.0576
[34/60][50/79]	Loss_D: 0.2115	Loss_G: 3.2222	D(x): 0.9098	D(G(z)): 0.0680 / 0.0538
[35/60][0/79]	Loss_D: 0.5568	Loss_G: 2.0906	D(x): 0.9084	D(G(z)): 0.3163 / 0.1545
[35/60][50/79]	Loss_D: 5.1111	Loss_G: 2.8137	D(x): 0.0195	D(G(z)): 0.0003 / 0.1043
[36/60][0/79]	Loss_D: 0.2795	Loss_G: 4.4470	D(x): 0.9940	D(G(z)): 0.2042 / 0.0178
[36/60][50/79]	Loss_D: 0.1687	Loss_G: 3.6446	D(x): 0.8880	D(G(z)): 0.0350 / 0.0466
[37/60][0/79]	Loss_D: 0.4623	Loss_G: 2.6764	D(x): 0.7381	D(G(z)): 0.0827 / 0.1034
[37/60][50/79]	Loss_D: 0.2594	Loss_G: 4.6265	D(x): 0.7939	D(G(z)): 0.0019 / 0.0156
[38/60][0/79]	Loss_D: 0.0417	Loss_G: 4.1340	D(x): 0.9848	D(G(z)): 0.0257 / 0.0207
[38/60][50/79]	Loss_D: 0.0613	Loss_G: 4.7990	D(x): 0.9868	D(G(z)): 0.0423 / 0.0128
[39/60][0/79]	Loss_D: 0.0126	Loss_G: 4.5422	D(x): 0.9990	D(G(z)): 0.0114 / 0.0161
[39/60][50/79]	Loss_D: 0.5382	Loss_G: 2.0068	D(x): 0.6838	D(G(z)): 0.0692 / 0.1815
[40/60][0/79]	Loss_D: 0.1843	Loss_G: 4.9301	D(x): 0.9713	D(G(z)): 0.1123 / 0.0147
[40/60][50/79]	Loss_D: 0.0114	Loss_G: 6.9595	D(x): 0.9983	D(G(z)): 0.0095 / 0.0026
[41/60][0/79]	Loss_D: 2.7421	Loss_G: 5.6773	D(x): 0.9543	D(G(z)): 0.8748 / 0.0198
[41/60][50/79]	Loss_D: 0.1682	Loss_G: 5.2652	D(x): 0.8612	D(G(z)): 0.0027 / 0.0082
[42/60][0/79]	Loss_D: 0.1054	Loss_G: 2.9872	D(x): 0.9965	D(G(z)): 0.0946 / 0.0626
[42/60][50/79]	Loss_D: 0.2524	Loss_G: 3.1311	D(x): 0.8583	D(G(z)): 0.0189 / 0.0583
[43/60][0/79]	Loss_D: 4.2940	Loss_G: 9.1384	D(x): 0.9991	D(G(z)): 0.9634 / 0.0003
[43/60][50/79]	Loss_D: 0.9087	Loss_G: 4.4272	D(x): 0.9725	D(G(z)): 0.4980 / 0.0178
[44/60][0/79]	Loss_D: 0.2963	Loss_G: 3.1727	D(x): 0.9064	D(G(z)): 0.1285 / 0.0696
[44/60][50/79]	Loss_D: 0.2790	Loss_G: 3.5738	D(x): 0.8128	D(G(z)): 0.0558 / 0.0438
[45/60][0/79]	Loss_D: 0.1481	Loss_G: 3.9010	D(x): 0.9132	D(G(z)): 0.0422 / 0.0468
[45/60][50/79]	Loss_D: 0.0852	Loss_G: 4.6652	D(x): 0.9938	D(G(z)): 0.0742 / 0.0132
[46/60][0/79]	Loss_D: 1.6991	Loss_G: 8.7159	D(x): 0.9984	D(G(z)): 0.7377 / 0.0004
[46/60][50/79]	Loss_D: 0.8390	Loss_G: 4.4246	D(x): 0.9862	D(G(z)): 0.5092 / 0.0215
[47/60][0/79]	Loss_D: 4.9910	Loss_G: 4.1055	D(x): 0.9999	D(G(z)): 0.9751 / 0.0324
[47/60][50/79]	Loss_D: 0.6508	Loss_G: 3.2114	D(x): 0.9895	D(G(z)): 0.4469 / 0.0477
[48/60][0/79]	Loss_D: 2.8913	Loss_G: 7.9553	D(x): 0.9998	D(G(z)): 0.9029 / 0.0007
[48/60][50/79]	Loss_D: 0.1558	Loss_G: 3.5601	D(x): 0.9295	D(G(z)): 0.0748 / 0.0406
[49/60][0/79]	Loss_D: 0.5910	Loss_G: 4.2077	D(x): 1.0000	D(G(z)): 0.3580 / 0.0241

[49/60][50/79]	Loss_D: 0.1235	Loss_G: 3.4171	$D(x)$: 0.9291	$D(G(z))$: 0.0383 / 0.0412
[50/60][0/79]	Loss_D: 4.3362	Loss_G: 3.9284	$D(x)$: 0.9998	$D(G(z))$: 0.9172 / 0.0926
[50/60][50/79]	Loss_D: 0.1021	Loss_G: 4.4575	$D(x)$: 0.9825	$D(G(z))$: 0.0763 / 0.0172
[51/60][0/79]	Loss_D: 0.6642	Loss_G: 3.2043	$D(x)$: 0.8651	$D(G(z))$: 0.3584 / 0.0626
[51/60][50/79]	Loss_D: 2.3933	Loss_G: 0.7965	$D(x)$: 0.1250	$D(G(z))$: 0.0004 / 0.5122
[52/60][0/79]	Loss_D: 0.1220	Loss_G: 4.1970	$D(x)$: 0.9808	$D(G(z))$: 0.0951 / 0.0202
[52/60][50/79]	Loss_D: 0.1891	Loss_G: 8.3906	$D(x)$: 0.8488	$D(G(z))$: 0.0001 / 0.0005
[53/60][0/79]	Loss_D: 0.0065	Loss_G: 7.2400	$D(x)$: 0.9949	$D(G(z))$: 0.0013 / 0.0012
[53/60][50/79]	Loss_D: 0.0918	Loss_G: 4.4401	$D(x)$: 0.9776	$D(G(z))$: 0.0592 / 0.0192
[54/60][0/79]	Loss_D: 4.1998	Loss_G: 6.7730	$D(x)$: 1.0000	$D(G(z))$: 0.8662 / 0.0017
[54/60][50/79]	Loss_D: 0.0901	Loss_G: 5.4064	$D(x)$: 0.9909	$D(G(z))$: 0.0679 / 0.0081
[55/60][0/79]	Loss_D: 0.3984	Loss_G: 5.8680	$D(x)$: 0.9924	$D(G(z))$: 0.2578 / 0.0041
[55/60][50/79]	Loss_D: 0.3112	Loss_G: 4.1508	$D(x)$: 0.7506	$D(G(z))$: 0.0031 / 0.0313
[56/60][0/79]	Loss_D: 2.0480	Loss_G: 8.0220	$D(x)$: 0.9211	$D(G(z))$: 0.7479 / 0.0008
[56/60][50/79]	Loss_D: 0.2949	Loss_G: 3.0195	$D(x)$: 0.9773	$D(G(z))$: 0.2043 / 0.0651
[57/60][0/79]	Loss_D: 0.1841	Loss_G: 3.6742	$D(x)$: 0.9251	$D(G(z))$: 0.0599 / 0.0691
[57/60][50/79]	Loss_D: 0.1021	Loss_G: 4.3927	$D(x)$: 0.9642	$D(G(z))$: 0.0615 / 0.0182
[58/60][0/79]	Loss_D: 0.0674	Loss_G: 4.5956	$D(x)$: 0.9951	$D(G(z))$: 0.0568 / 0.0175
[58/60][50/79]	Loss_D: 0.0491	Loss_G: 5.4625	$D(x)$: 0.9969	$D(G(z))$: 0.0440 / 0.0067
[59/60][0/79]	Loss_D: 2.1525	Loss_G: 7.6566	$D(x)$: 0.9996	$D(G(z))$: 0.7589 / 0.0012
[59/60][50/79]	Loss_D: 0.0310	Loss_G: 4.2983	$D(x)$: 0.9979	$D(G(z))$: 0.0280 / 0.0181

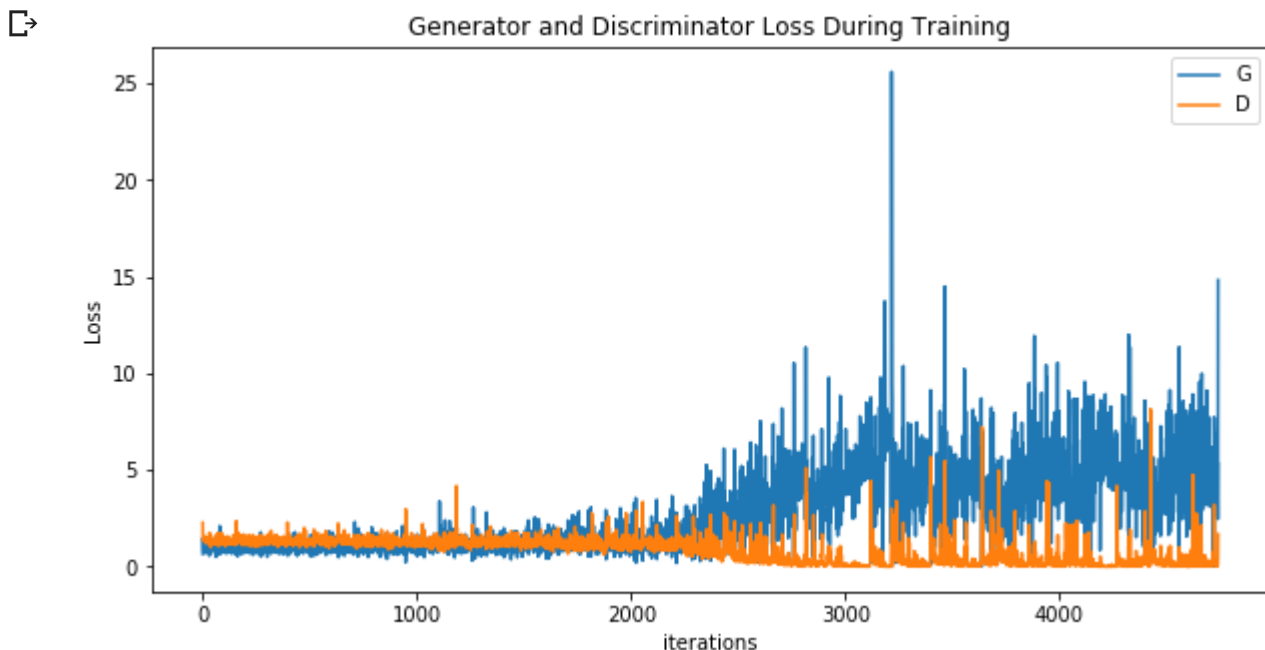
▼ Results

Finally, let's check out how we did. Here, we will look at three different results. First, we will see how D changed during training. Second, we will visualize G's output on the fixed_noise batch for every epoch. We will look at a batch of real data next to a batch of fake data from G.

Loss versus training iteration

Below is a plot of D & G's losses versus training iterations.

```
plt.figure(figsize=(10,5))
plt.title("Generator and Discriminator Loss During Training")
plt.plot(G_losses, label="G")
plt.plot(D_losses, label="D")
plt.xlabel("iterations")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



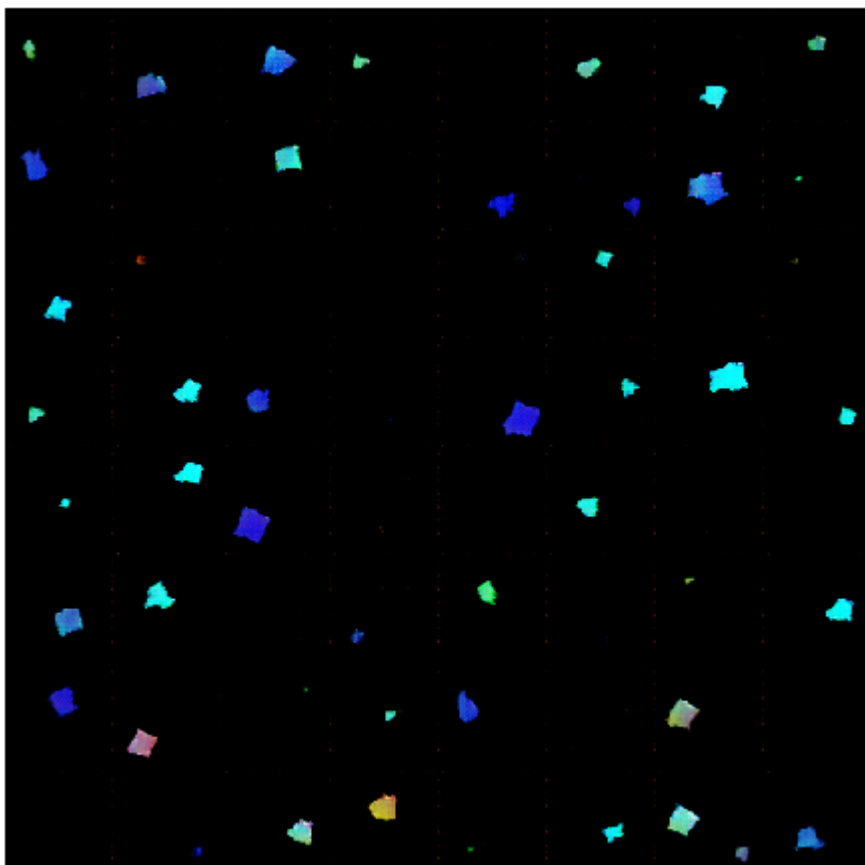
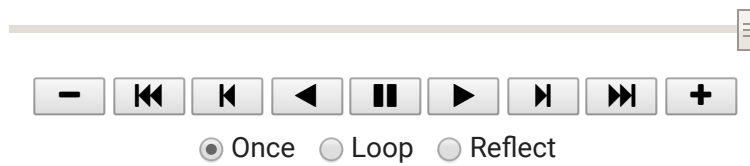
Visualization of G's progression

Remember how we saved the generator's output on the fixed_noise batch after every epoch of training. We can now visualize the training progression of G with an animation. Press the play button to start the animation.

```
%%capture
fig = plt.figure(figsize=(8,8))
plt.axis("off")
ims = [[plt.imshow(np.transpose(i,(1,2,0)), animated=True)] for i in img_list]
ani = animation.ArtistAnimation(fig, ims, interval=1000, repeat_delay=1000, blit=True)

HTML(ani.to_jshtml())
```





Real Images vs. Fake Images

Finally, lets take a look at some real images and fake images side by side.

```
# Grab a batch of real images from the dataloader
real_batch = next(iter(dataloader))

# Plot the real images
plt.figure(figsize=(15,15))
plt.subplot(1,2,1)
plt.axis("off")
plt.title("Real Images")
plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:64], padding=5, normalize=True).data, [0, 2, 1]))

# Plot the fake images from the last epoch
plt.subplot(1,2,2)
plt.axis("off")
plt.title("Fake Images")
plt.imshow(np.transpose(img_list[-1], (1, 2, 0)))
plt.show()
```

Where to Go Next

We have reached the end of our journey, but there are several places you could go from here. You could

- Train for longer to see how good the results get
- Modify this model to take a different dataset and possibly change the size of the images and the architecture
- Check out some other cool GAN projects here <<https://github.com/nashory/gans-awesome-apps>>
- Create GANs that generate music <<https://deepmind.com/blog/wavenet-generative-model-raises-the-bar-for-machine-generated-sounds/>>