

```
from google.colab import drive
drive.mount('/content/drive')
```

➞ Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=9473189

Enter your authorization code:

.....

Mounted at /content/drive

```
!unzip '/content/drive/My Drive/Flickr-dog.zip' -d '/content/sample_data/'
```

```
!apt-get install p7zip-full
```

```
➞ Reading package lists... Done
Building dependency tree
Reading state information... Done
p7zip-full is already the newest version (16.02+dfsg-6).
0 upgraded, 0 newly installed, 0 to remove and 28 not upgraded.
```

```
!tar -xvf DogsImages.tar -C /content/sample_data/
```

```
from __future__ import print_function
%matplotlib inline
import argparse
import os
import random
import torch
import torch.nn as nn
import torch.nn.parallel
import torch.backends.cudnn as cudnn
import torch.optim as optim
import torch.utils.data
import torchvision.datasets as dset
import torchvision.transforms as transforms
import torchvision.utils as vutils
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.animation as animation
from IPython.display import HTML
import zipfile
import cv2
# import unzip
```

```
# Set random seem for reproducibility
manualSeed = 999
```

```

#manualSeed = random.randint(1, 10000) # use if you want new results
print("Random Seed: ", manualSeed)
random.seed(manualSeed)
torch.manual_seed(manualSeed)

```

```

☞ Random Seed: 999
<torch._C.Generator at 0x7f71b7b4d0b0>

```

```

# drive = '/content/sample_data/square_data/squares/'
# for i in range(5000):
#     x = random.randrange(0,5)
#     img = np.zeros((64,64,3), np.uint8)
#     color = (random.randrange(1,255), random.randrange(1,255), random.randrange(1,255))

#     img[42-x:42+x, 42-x:42+x, :] = tuple(reversed(color))
#     rotationMat = cv2.getRotationMatrix2D((32,32), random.randrange(0,90), 1)
#     img = cv2.warpAffine(img, rotationMat,(img.shape[0],img.shape[1]))
#     trans_mat = np.float32([[1,0,random.randrange(-20,20)], [0,1,random.randrange(-20,20)]])
#     img = cv2.warpAffine(img, trans_mat, (img.shape[0],img.shape[1]))
#     # cv2.war
#     file = drive + 'square'+ str(i) + '.png'
#     plt.imsave(file, img)

```

▼ Inputs

Let's define some inputs for the run:

- **dataroot** - the path to the root of the dataset folder. We will talk more about the dataset in the next section.
- **workers** - the number of worker threads for loading the data with the DataLoader
- **batch_size** - the batch size used in training. The DCGAN paper uses a batch size of 128
- **image_size** - the spatial size of the images used for training. This implementation defaults to 64x64. If you want to use larger images, the structures of D and G must be changed. See [here](https://github.com/pytorch/examples/issues/100) <<https://github.com/pytorch/examples/issues/100>>
- **nc** - number of color channels in the input images. For color images this is 3
- **nz** - length of latent vector
- **ngf** - relates to the depth of feature maps carried through the generator
- **ndf** - sets the depth of feature maps propagated through the discriminator
- **num_epochs** - number of training epochs to run. Training for longer will probably lead to better results
- **lr** - learning rate for training. As described in the DCGAN paper, this number should be 0.0002
- **beta1** - beta1 hyperparameter for Adam optimizers. As described in paper, this number should be 0.5
- **ngpu** - number of GPUs available. If this is 0, code will run in CPU mode. If this number is greater than 0, it will use that many GPUs.

```

# Root directory for dataset

```

```
dataroot = "/content/sample_data/Flickr-dog"

# Number of workers for dataloader
workers = 4

# Batch size during training
batch_size = 32

# Spatial size of training images. All images will be resized to this
# size using a transformer.
image_size = 64

# Number of channels in the training images. For color images this is 3
nc = 3

# Size of z latent vector (i.e. size of generator input)
nz = 100

# Size of feature maps in generator
ngf = 64

# Size of feature maps in discriminator
ndf = 64

# Number of training epochs
num_epochs = 100

# Learning rate for optimizers
lr = 0.0005

# Beta1 hyperparam for Adam optimizers
beta1 = 0.5

# Number of GPUs available. Use 0 for CPU mode.
ngpu = 1

# We can use an image folder dataset the way we have it setup.
# Create the dataset
dataset = dset.ImageFolder(root=dataroot,
                           transform=transforms.Compose([
                               transforms.Resize(image_size),
                               transforms.CenterCrop(image_size),
                               transforms.ToTensor(),
                               transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
                           ]))

# Create the dataloader
dataloader = torch.utils.data.DataLoader(dataset, batch_size=batch_size,
                                         shuffle=True, num_workers=workers)

# Decide which device we want to run on
device = torch.device("cuda:0" if (torch.cuda.is_available() and ngpu > 0) else "cpu")
```

```
# Plot some training images
real_batch = next(iter(dataloader))
plt.figure(figsize=(8,8))
plt.axis("off")
plt.title("Training Images")
plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:64], padding=2, normalize=
```

↳ <matplotlib.image.AxesImage at 0x7f71ab5702b0>



```
# custom weights initialization called on netG and netD
```

```
def weights_init(m):
    classname = m.__class__.__name__
    if classname.find('Conv') != -1:
        nn.init.normal_(m.weight.data, 0.0, 0.02)
    elif classname.find('BatchNorm') != -1:
        nn.init.normal_(m.weight.data, 1.0, 0.02)
        nn.init.constant_(m.bias.data, 0)
```

```
# Generator Code
```

```
class Generator(nn.Module):
    def __init__(self, ngpu):
        super(Generator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is Z, going into a convolution
            nn.ConvTranspose2d( nz, ngf * 8, 4, 1, 0, bias=False),
            nn.BatchNorm2d(ngf * 8),
            nn.ReLU(True),
            # state size. (ngf*8) x 4 x 4
            nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True),
            # state size. (ngf*4) x 8 x 8
```

```

nn.ConvTranspose2d( ngf * 4, ngf * 2, 4, 2, 1, bias=False),
nn.BatchNorm2d(ngf * 2),
nn.ReLU(True),
# state size. (ngf*2) x 16 x 16
nn.ConvTranspose2d( ngf * 2, ngf, 4, 2, 1, bias=False),
nn.BatchNorm2d(ngf),
nn.ReLU(True),
# state size. (ngf) x 32 x 32
nn.ConvTranspose2d( ngf, nc, 4, 2, 1, bias=False),
nn.Tanh()
# state size. (nc) x 64 x 64
)

def forward(self, input):
    return self.main(input)

```

Now, we can instantiate the generator and apply the `weights_init` function. Check out the printed model structure.

```

# Create the generator
netG = Generator(ngpu).to(device)

# Handle multi-gpu if desired
if (device.type == 'cuda') and (ngpu > 1):
    netG = nn.DataParallel(netG, list(range(ngpu)))

# Apply the weights_init function to randomly initialize all weights
# to mean=0, stdev=0.2.
netG.apply(weights_init)

# Print the model
print(netG)

↳ Generator(
  (main): Sequential(
    (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU(inplace=True)
    (9): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (11): ReLU(inplace=True)
    (12): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (13): Tanh()
  )
)

```

Discriminator Code

```
class Discriminator(nn.Module):
    def __init__(self, ngpu):
        super(Discriminator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is (nc) x 64 x 64
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf) x 32 x 32
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 2),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*2) x 16 x 16
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 4),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*4) x 8 x 8
            nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 8),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*8) x 4 x 4
            nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),
            nn.Sigmoid()
        )

    def forward(self, input):
        return self.main(input)
```

Now, as with the generator, we can create the discriminator, apply the `weights_init` function, and pri

```
# Create the Discriminator
netD = Discriminator(ngpu).to(device)

# Handle multi-gpu if desired
if (device.type == 'cuda') and (ngpu > 1):
    netD = nn.DataParallel(netD, list(range(ngpu)))

# Apply the weights_init function to randomly initialize all weights
# to mean=0, stdev=0.2.
netD.apply(weights_init)

# Print the model
print(netD)
```



```

Discriminator(
  (main): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (10): LeakyReLU(negative_slope=0.2, inplace=True)
    (11): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (12): Sigmoid()
  )
)

```

```
# Initialize BCELoss function
```

```
criterion = nn.BCELoss()
```

```
# Create batch of latent vectors that we will use to visualize
```

```
# the progression of the generator
```

```
fixed_noise = torch.randn(64, nz, 1, 1, device=device)
```

```
# Establish convention for real and fake labels during training
```

```
real_label = 1
```

```
fake_label = 0
```

```
# Setup Adam optimizers for both G and D
```

```
optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
```

```
optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))
```

```
# Training Loop
```

```
# Lists to keep track of progress
```

```
img_list = []
```

```
G_losses = []
```

```
D_losses = []
```

```
iters = 0
```

```
print("Starting Training Loop...")
```

```
# For each epoch
```

```
for epoch in range(num_epochs):
```

```
    # For each batch in the dataloader
```

```
    for i, data in enumerate(dataloader, 0):
```

```
        #####
```

```
        # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
```

```
        #####
```

```
        ## Train with all-real batch
```

```

netD.zero_grad()
# Format batch
real_cpu = data[0].to(device)
b_size = real_cpu.size(0)
label = torch.full((b_size,), real_label, device=device)
# Forward pass real batch through D
output = netD(real_cpu).view(-1)
# Calculate loss on all-real batch
errD_real = criterion(output, label)
# Calculate gradients for D in backward pass
errD_real.backward()
D_x = output.mean().item()

## Train with all-fake batch
# Generate batch of latent vectors
noise = torch.randn(b_size, nz, 1, 1, device=device)
# Generate fake image batch with G
fake = netG(noise)
label.fill_(fake_label)
# Classify all fake batch with D
output = netD(fake.detach()).view(-1)
# Calculate D's loss on the all-fake batch
errD_fake = criterion(output, label)
# Calculate the gradients for this batch
errD_fake.backward()
D_G_z1 = output.mean().item()
# Add the gradients from the all-real and all-fake batches
errD = errD_real + errD_fake
# Update D
optimizerD.step()

#####
# (2) Update G network: maximize log(D(G(z)))
#####
netG.zero_grad()
label.fill_(real_label) #fake labels are real for generator cost
# Since we just updated D, perform another forward pass of all-fake batch through D
output = netD(fake).view(-1)
# Calculate G's loss based on this output
errG = criterion(output, label)
# Calculate gradients for G
errG.backward()
D_G_z2 = output.mean().item()
# Update G
optimizerG.step()

# Output training stats
if i % 50 == 0:
    print('%d/%d [%d/%d] \tLoss_D: %.4f \tLoss_G: %.4f \tD(x): %.4f \tD(G(z)): %.4f / %.
          % (epoch, num_epochs, i, len(dataloader),
              errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))

```



```
# Save Losses for plotting later
G_losses.append(errG.item())
D_losses.append(errD.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 100 == 0) or ((epoch == num_epochs-1) and (i == len(dataloader)-1)):
    with torch.no_grad():
        fake = netG(fixed_noise).detach().cpu()
        img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

iters += 1
```



Starting Training Loop...

[0/100][0/12]	Loss_D: 0.4606	Loss_G: 7.8754	D(x): 0.9659	D(G(z)): 0.2564 / 0.0009
[1/100][0/12]	Loss_D: 0.1862	Loss_G: 4.5960	D(x): 0.9394	D(G(z)): 0.0992 / 0.0185
[2/100][0/12]	Loss_D: 0.2264	Loss_G: 4.9598	D(x): 0.9258	D(G(z)): 0.1178 / 0.0131
[3/100][0/12]	Loss_D: 0.2165	Loss_G: 5.4288	D(x): 0.9006	D(G(z)): 0.0883 / 0.0105
[4/100][0/12]	Loss_D: 0.2408	Loss_G: 5.7874	D(x): 0.9047	D(G(z)): 0.0839 / 0.0071
[5/100][0/12]	Loss_D: 0.1308	Loss_G: 5.3680	D(x): 0.9438	D(G(z)): 0.0644 / 0.0115
[6/100][0/12]	Loss_D: 0.4179	Loss_G: 8.1568	D(x): 0.9725	D(G(z)): 0.2502 / 0.0007
[7/100][0/12]	Loss_D: 0.6443	Loss_G: 5.4085	D(x): 0.7318	D(G(z)): 0.1180 / 0.0136
[8/100][0/12]	Loss_D: 0.3162	Loss_G: 5.6046	D(x): 0.9264	D(G(z)): 0.1086 / 0.0074
[9/100][0/12]	Loss_D: 0.1338	Loss_G: 5.2713	D(x): 0.9542	D(G(z)): 0.0676 / 0.0095
[10/100][0/12]	Loss_D: 0.4878	Loss_G: 8.0307	D(x): 0.9371	D(G(z)): 0.1855 / 0.0062
[11/100][0/12]	Loss_D: 0.2537	Loss_G: 5.4491	D(x): 0.9483	D(G(z)): 0.1642 / 0.0129
[12/100][0/12]	Loss_D: 0.1785	Loss_G: 5.3023	D(x): 0.9253	D(G(z)): 0.0688 / 0.0164
[13/100][0/12]	Loss_D: 0.1382	Loss_G: 5.9818	D(x): 0.9234	D(G(z)): 0.0445 / 0.0057
[14/100][0/12]	Loss_D: 0.1763	Loss_G: 5.0628	D(x): 0.9072	D(G(z)): 0.0539 / 0.0135
[15/100][0/12]	Loss_D: 0.2831	Loss_G: 5.3494	D(x): 0.8712	D(G(z)): 0.0888 / 0.0117
[16/100][0/12]	Loss_D: 0.7262	Loss_G: 10.4547	D(x): 0.9827	D(G(z)): 0.3724 / 0.0001
[17/100][0/12]	Loss_D: 0.3887	Loss_G: 5.2466	D(x): 0.8774	D(G(z)): 0.1070 / 0.0226
[18/100][0/12]	Loss_D: 0.4538	Loss_G: 8.5613	D(x): 0.9486	D(G(z)): 0.2296 / 0.0014
[19/100][0/12]	Loss_D: 0.1833	Loss_G: 7.4461	D(x): 0.9591	D(G(z)): 0.0969 / 0.0018
[20/100][0/12]	Loss_D: 0.1336	Loss_G: 5.1001	D(x): 0.9409	D(G(z)): 0.0578 / 0.0111
[21/100][0/12]	Loss_D: 0.0978	Loss_G: 6.0094	D(x): 0.9593	D(G(z)): 0.0493 / 0.0053
[22/100][0/12]	Loss_D: 0.2283	Loss_G: 7.3405	D(x): 0.9914	D(G(z)): 0.1518 / 0.0009
[23/100][0/12]	Loss_D: 0.1869	Loss_G: 6.2461	D(x): 0.9394	D(G(z)): 0.1026 / 0.0039
[24/100][0/12]	Loss_D: 0.0542	Loss_G: 5.7488	D(x): 0.9682	D(G(z)): 0.0199 / 0.0066
[25/100][0/12]	Loss_D: 0.1576	Loss_G: 5.8005	D(x): 0.9781	D(G(z)): 0.1121 / 0.0058
[26/100][0/12]	Loss_D: 0.1305	Loss_G: 5.6543	D(x): 0.9553	D(G(z)): 0.0691 / 0.0062
[27/100][0/12]	Loss_D: 0.2468	Loss_G: 7.0703	D(x): 0.9923	D(G(z)): 0.1539 / 0.0036
[28/100][0/12]	Loss_D: 0.1727	Loss_G: 6.2242	D(x): 0.8831	D(G(z)): 0.0284 / 0.0099
[29/100][0/12]	Loss_D: 0.3905	Loss_G: 7.0902	D(x): 0.8420	D(G(z)): 0.0184 / 0.0238
[30/100][0/12]	Loss_D: 0.8110	Loss_G: 10.5747	D(x): 0.9417	D(G(z)): 0.3813 / 0.0007
[31/100][0/12]	Loss_D: 0.4732	Loss_G: 10.2280	D(x): 0.9777	D(G(z)): 0.2924 / 0.0012
[32/100][0/12]	Loss_D: 0.2470	Loss_G: 7.3577	D(x): 0.9327	D(G(z)): 0.1172 / 0.0018
[33/100][0/12]	Loss_D: 0.2783	Loss_G: 8.9674	D(x): 0.9152	D(G(z)): 0.1362 / 0.0025
[34/100][0/12]	Loss_D: 0.1655	Loss_G: 6.5588	D(x): 0.9478	D(G(z)): 0.0800 / 0.0045
[35/100][0/12]	Loss_D: 0.1825	Loss_G: 6.2660	D(x): 0.9869	D(G(z)): 0.1311 / 0.0068
[36/100][0/12]	Loss_D: 0.4751	Loss_G: 11.4205	D(x): 0.9865	D(G(z)): 0.2758 / 0.0004
[37/100][0/12]	Loss_D: 0.1478	Loss_G: 5.9893	D(x): 0.9504	D(G(z)): 0.0700 / 0.0059
[38/100][0/12]	Loss_D: 0.1601	Loss_G: 7.9928	D(x): 0.9878	D(G(z)): 0.1238 / 0.0016
[39/100][0/12]	Loss_D: 0.0710	Loss_G: 6.7695	D(x): 0.9727	D(G(z)): 0.0400 / 0.0039
[40/100][0/12]	Loss_D: 0.0584	Loss_G: 6.2230	D(x): 0.9856	D(G(z)): 0.0409 / 0.0045
[41/100][0/12]	Loss_D: 0.0463	Loss_G: 5.8976	D(x): 0.9792	D(G(z)): 0.0236 / 0.0068
[42/100][0/12]	Loss_D: 0.0396	Loss_G: 6.2987	D(x): 0.9795	D(G(z)): 0.0180 / 0.0046
[43/100][0/12]	Loss_D: 0.0446	Loss_G: 8.3428	D(x): 0.9674	D(G(z)): 0.0093 / 0.0017
[44/100][0/12]	Loss_D: 0.0683	Loss_G: 6.2304	D(x): 0.9888	D(G(z)): 0.0475 / 0.0042
[45/100][0/12]	Loss_D: 0.0580	Loss_G: 6.6429	D(x): 0.9860	D(G(z)): 0.0404 / 0.0044
[46/100][0/12]	Loss_D: 0.0932	Loss_G: 6.5549	D(x): 0.9841	D(G(z)): 0.0680 / 0.0031
[47/100][0/12]	Loss_D: 0.6402	Loss_G: 10.1137	D(x): 0.9381	D(G(z)): 0.2609 / 0.0013
[48/100][0/12]	Loss_D: 0.3478	Loss_G: 9.6045	D(x): 0.9248	D(G(z)): 0.1520 / 0.0016
[49/100][0/12]	Loss_D: 1.1717	Loss_G: 7.3010	D(x): 0.6630	D(G(z)): 0.1406 / 0.0204
[50/100][0/12]	Loss_D: 0.4913	Loss_G: 9.2208	D(x): 0.7613	D(G(z)): 0.0229 / 0.0415
[51/100][0/12]	Loss_D: 1.1828	Loss_G: 7.0112	D(x): 0.4549	D(G(z)): 0.0048 / 0.0451
[52/100][0/12]	Loss_D: 0.2182	Loss_G: 7.2926	D(x): 0.9361	D(G(z)): 0.1070 / 0.0063
[53/100][0/12]	Loss_D: 0.1773	Loss_G: 7.4166	D(x): 0.8988	D(G(z)): 0.0156 / 0.0048
[54/100][0/12]	Loss_D: 0.1286	Loss_G: 6.2247	D(x): 0.9754	D(G(z)): 0.0864 / 0.0039
[55/100][0/12]	Loss_D: 0.0595	Loss_G: 5.3710	D(x): 0.9690	D(G(z)): 0.0241 / 0.0091

```

[56/100][0/12] Loss_D: 0.0567 Loss_G: 6.3720 D(x): 0.9758 D(G(z)): 0.0298 / 0.0041
[57/100][0/12] Loss_D: 0.1755 Loss_G: 8.0688 D(x): 0.9960 D(G(z)): 0.1343 / 0.0005
[58/100][0/12] Loss_D: 0.0416 Loss_G: 6.2708 D(x): 0.9742 D(G(z)): 0.0144 / 0.0046
[59/100][0/12] Loss_D: 0.1344 Loss_G: 8.8964 D(x): 0.9893 D(G(z)): 0.1003 / 0.0013
[60/100][0/12] Loss_D: 0.1411 Loss_G: 7.2474 D(x): 0.9821 D(G(z)): 0.0655 / 0.0085
[61/100][0/12] Loss_D: 0.0497 Loss_G: 5.1717 D(x): 0.9716 D(G(z)): 0.0183 / 0.0129
[62/100][0/12] Loss_D: 0.0815 Loss_G: 6.1618 D(x): 0.9904 D(G(z)): 0.0606 / 0.0037
[63/100][0/12] Loss_D: 0.0411 Loss_G: 6.3632 D(x): 0.9831 D(G(z)): 0.0227 / 0.0052
[64/100][0/12] Loss_D: 0.0206 Loss_G: 6.2924 D(x): 0.9909 D(G(z)): 0.0112 / 0.0039
[65/100][0/12] Loss_D: 0.0843 Loss_G: 7.3455 D(x): 0.9899 D(G(z)): 0.0595 / 0.0013
[66/100][0/12] Loss_D: 0.0479 Loss_G: 5.8734 D(x): 0.9886 D(G(z)): 0.0346 / 0.0048
[67/100][0/12] Loss_D: 0.0868 Loss_G: 6.5758 D(x): 0.9943 D(G(z)): 0.0719 / 0.0022
[68/100][0/12] Loss_D: 0.0796 Loss_G: 6.8328 D(x): 0.9972 D(G(z)): 0.0679 / 0.0026
[69/100][0/12] Loss_D: 0.0514 Loss_G: 5.2122 D(x): 0.9711 D(G(z)): 0.0201 / 0.0113
[70/100][0/12] Loss_D: 0.0253 Loss_G: 5.4775 D(x): 0.9871 D(G(z)): 0.0119 / 0.0087
[71/100][0/12] Loss_D: 0.0325 Loss_G: 6.1019 D(x): 0.9913 D(G(z)): 0.0219 / 0.0051
[72/100][0/12] Loss_D: 0.0318 Loss_G: 6.5252 D(x): 0.9810 D(G(z)): 0.0118 / 0.0035
[73/100][0/12] Loss_D: 0.0251 Loss_G: 6.0494 D(x): 0.9939 D(G(z)): 0.0184 / 0.0062
[74/100][0/12] Loss_D: 0.0340 Loss_G: 6.8148 D(x): 0.9896 D(G(z)): 0.0217 / 0.0033
[75/100][0/12] Loss_D: 0.0384 Loss_G: 6.5937 D(x): 0.9930 D(G(z)): 0.0291 / 0.0041
[76/100][0/12] Loss_D: 0.0471 Loss_G: 7.7230 D(x): 0.9572 D(G(z)): 0.0017 / 0.0014
[77/100][0/12] Loss_D: 0.0472 Loss_G: 6.9048 D(x): 0.9682 D(G(z)): 0.0119 / 0.0027
[78/100][0/12] Loss_D: 0.0449 Loss_G: 6.3658 D(x): 0.9632 D(G(z)): 0.0063 / 0.0075
[79/100][0/12] Loss_D: 0.0229 Loss_G: 5.7054 D(x): 0.9893 D(G(z)): 0.0117 / 0.0064
[80/100][0/12] Loss_D: 0.1462 Loss_G: 8.9476 D(x): 0.9965 D(G(z)): 0.1089 / 0.0005
[81/100][0/12] Loss_D: 0.0475 Loss_G: 5.7444 D(x): 0.9613 D(G(z)): 0.0069 / 0.0134
[82/100][0/12] Loss_D: 0.0281 Loss_G: 5.8606 D(x): 0.9869 D(G(z)): 0.0144 / 0.0066
[83/100][0/12] Loss_D: 0.0441 Loss_G: 6.7813 D(x): 0.9865 D(G(z)): 0.0285 / 0.0029
[84/100][0/12] Loss_D: 0.0424 Loss_G: 6.1440 D(x): 0.9828 D(G(z)): 0.0236 / 0.0045
[85/100][0/12] Loss_D: 0.0786 Loss_G: 6.3675 D(x): 0.9342 D(G(z)): 0.0028 / 0.0046
[86/100][0/12] Loss_D: 0.1547 Loss_G: 6.4302 D(x): 0.8837 D(G(z)): 0.0007 / 0.0089
[87/100][0/12] Loss_D: 4.7125 Loss_G: 0.8092 D(x): 0.4666 D(G(z)): 0.5499 / 0.5697
[88/100][0/12] Loss_D: 1.7826 Loss_G: 2.0511 D(x): 0.5723 D(G(z)): 0.4396 / 0.2956
[89/100][0/12] Loss_D: 1.7346 Loss_G: 1.9253 D(x): 0.7037 D(G(z)): 0.5625 / 0.2266
[90/100][0/12] Loss_D: 1.1681 Loss_G: 1.6703 D(x): 0.7102 D(G(z)): 0.4111 / 0.2534
[91/100][0/12] Loss_D: 1.2103 Loss_G: 2.3459 D(x): 0.7785 D(G(z)): 0.5191 / 0.1521
[92/100][0/12] Loss_D: 0.9493 Loss_G: 2.1241 D(x): 0.5836 D(G(z)): 0.1954 / 0.1696
[93/100][0/12] Loss_D: 1.6361 Loss_G: 2.6665 D(x): 0.7286 D(G(z)): 0.5965 / 0.1215
[94/100][0/12] Loss_D: 1.2198 Loss_G: 2.6541 D(x): 0.6286 D(G(z)): 0.4099 / 0.1223
[95/100][0/12] Loss_D: 1.3874 Loss_G: 3.7668 D(x): 0.7326 D(G(z)): 0.4602 / 0.1146
[96/100][0/12] Loss_D: 1.0122 Loss_G: 3.9321 D(x): 0.8461 D(G(z)): 0.4158 / 0.0649
[97/100][0/12] Loss_D: 0.9326 Loss_G: 3.2308 D(x): 0.7794 D(G(z)): 0.3694 / 0.0754
[98/100][0/12] Loss_D: 1.3985 Loss_G: 5.5441 D(x): 0.7832 D(G(z)): 0.5211 / 0.0176
[99/100][0/12] Loss_D: 0.4006 Loss_G: 4.6075 D(x): 0.9030 D(G(z)): 0.2066 / 0.0328

```

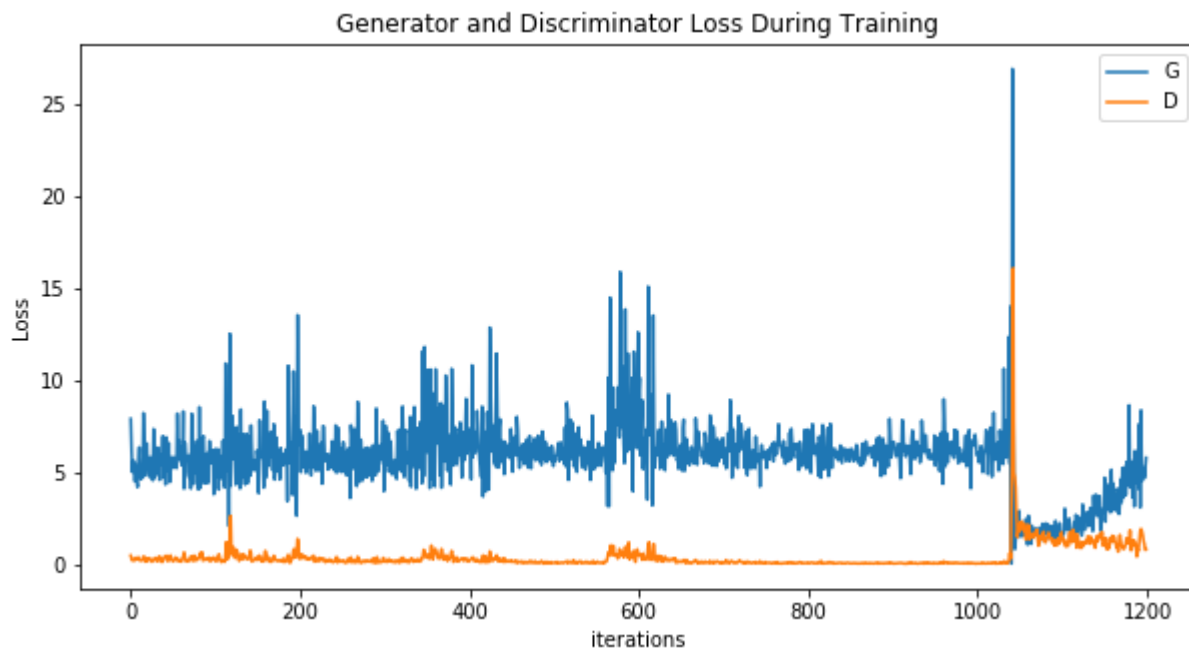
▼ Results

Finally, lets check out how we did. Here, we will look at three different results. First, we will see how D Second, we will visualize G's output on the fixed_noise batch for every epoch. And third, we will look at fake data from G.

Loss versus training iteration

Below is a plot of D & G's losses versus training iterations.

```
plt.figure(figsize=(10,5))
plt.title("Generator and Discriminator Loss During Training")
plt.plot(G_losses,label="G")
plt.plot(D_losses,label="D")
plt.xlabel("iterations")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



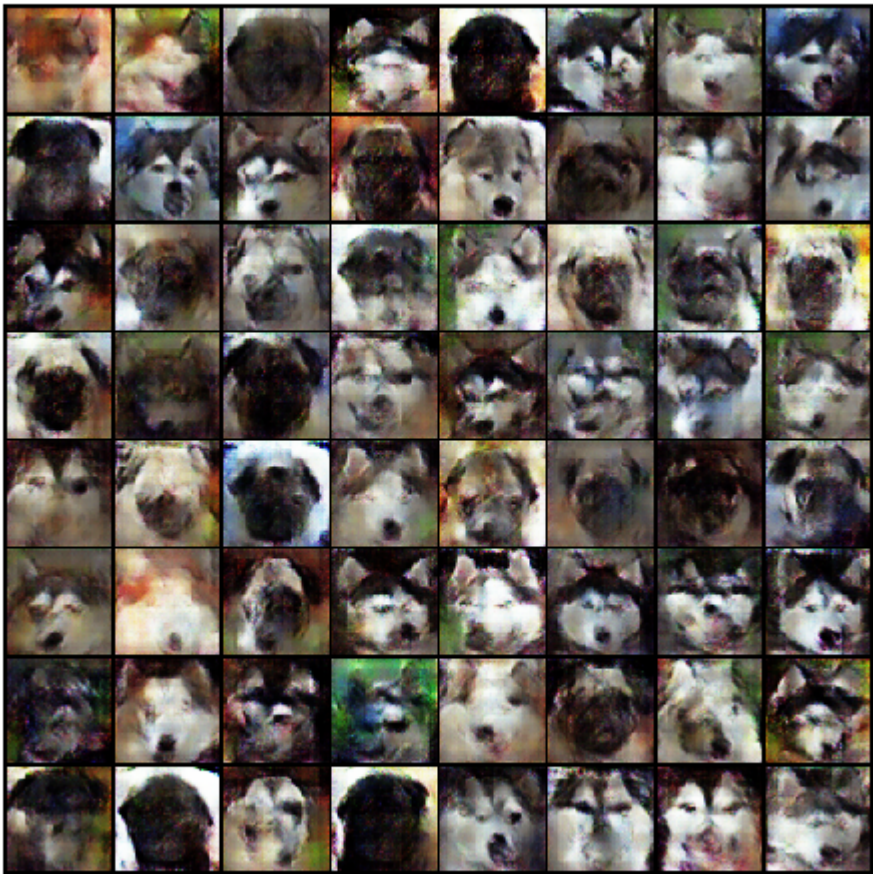
Visualization of G's progression

Remember how we saved the generator's output on the fixed_noise batch after every epoch of training progression of G with an animation. Press the play button to start the animation.

```
%%capture
fig = plt.figure(figsize=(8,8))
plt.axis("off")
ims = [[plt.imshow(np.transpose(i,(1,2,0))), animated=True]] for i in img_list]
ani = animation.ArtistAnimation(fig, ims, interval=1000, repeat_delay=1000, blit=True)

HTML(ani.to_jshtml())
```

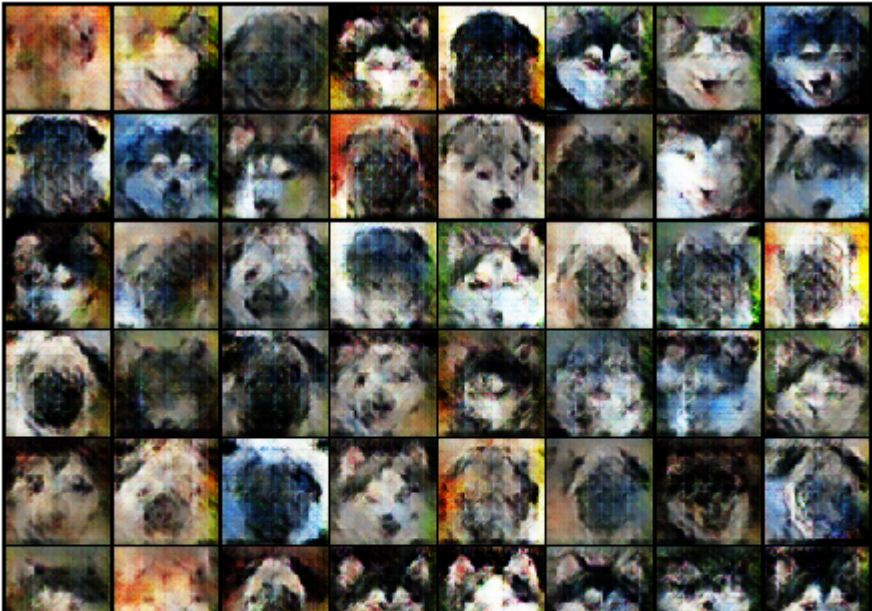


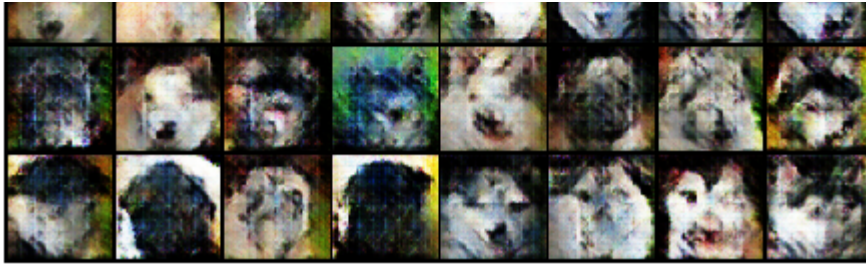


Slider control with a range from 0 to 100, currently set at approximately 60.

Navigation buttons: - (minus), << (double left arrow), < (single left arrow), || (pause), > (single right arrow), >> (double right arrow), + (plus).

Mode selection: ☐ Once ☒ Loop ☐ Reflect





Real Images vs. Fake Images

Finally, lets take a look at some real images and fake images side by side.

```
# Grab a batch of real images from the dataloader
real_batch = next(iter(dataloader))

# Plot the real images
plt.figure(figsize=(15,15))
plt.subplot(1,2,1)
plt.axis("off")
plt.title("Real Images")
plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:64], padding=5, normalize=

# Plot the fake images from the last epoch
plt.subplot(1,2,2)
plt.axis("off")
plt.title("Fake Images")
plt.imshow(np.transpose(img_list[-4],(1,2,0)))
plt.show()
```





▼ Where to Go Next

We have reached the end of our journey, but there are several places you could go from here. You could

- Train for longer to see how good the results get
- Modify this model to take a different dataset and possibly change the size of the images and the number of classes
- Check out some other cool GAN projects here <<https://github.com/nashory/gans-awesome-apps>>
- Create GANs that generate music <<https://deepmind.com/blog/wavenet-generative-model-raw-audio>>

